

Summary of Geigle’s “Inference Methods for LDA” (1)

Thomas W. Jones

6/25/2018

The Dirichlet Distribution and its Relation to the Multinomial

Before diving into LDA, it’s worth having some background on relevant probability distributions. The most relevant distributions are the categorical distribution, the multinomial distribution, and the Dirichlet distribution.

The categorical and multinomial distributions

First, the multinomial and categorical distributions are almost the same thing. The categorical distribution is like a Bernoulli distribution that has more than two outcomes. So, a K -dimensional categorical random variable will take on the value of 1 for one of its dimensions and 0 for the rest. The parameter of the categorical distribution is a K -dimensional vector of probabilities of each category. For example, a random variable representing a dingle role of a dice would be distributed categorical with a vector of probabilities equal to $[1/6, 1/6, 1/6, 1/6, 1/6, 1/6]$.

The multinomial distribution is a sum of categorical random variables, much like the binomial distribution is a sum of Bernoulli random variables. Instead of a K -dimensional vector of ones and zeros, a multinomial random variable is a vector of integers whose entries add up to n , the number of trials. Obviously, when $n = 1$ you have a categorical distribution.

The probability mass function of a multinomial random variable, \vec{y} with parameters n and $\vec{\theta}$ is below.

$$P(\vec{y}|\vec{\theta}) = \frac{n!}{\prod_{k=1}^K y_k!} \prod_{k=1}^K \theta_k^{y_k} \quad (1)$$

The multinomial distribution is often a good choice for modeling counts of words. In the case of LDA, our formulations have $n = 1$ due to the nature of the two-step sampling process it models. (Choose a topic, now given the topic, choose a word. Repeat.) In the LDA literature, the terms “categorical distribution” and “multinomial distribution” are often conflated. But because they are equivalent when $n = 1$, this (hopefully) does not cause much confusion.

The Dirichlet distribution

The Dirichlet distribution is the multivariate form of the beta distribution. Like the beta-binomial, the Dirichlet distribution is a conjugate prior for the multinomial. Conjugacy is established in the Bayesian context when the posterior and prior have the same functional form. Conjugacy is not required for Bayesian modeling. Yet using conjugate distributions tends to simplify the algebra and can be estimated in a more computationally-efficient manner.

If $\vec{\theta}$ is a K -dimensional random variable which is distributed Dirichlet, with parameter $\vec{\alpha}$, then its probability density function is given by

$$P(\vec{\theta}|\vec{\alpha}) = \mathcal{B}(\vec{\alpha})^{-1} \prod_{k=1}^K \theta_k^{\alpha_k-1} \quad (2)$$

Where $\mathcal{B}(\vec{\alpha})$ is the beta function and equal to $\frac{\prod_{k=1}^K \Gamma(\alpha_k)}{\Gamma(\sum_{k=1}^K \alpha_k)}$.

The Dirichlet distribution has a few properties that (according to Geigle) will come in handy later.

Proof of conjugacy WRT the multinomial

Here, we show that the Dirichlet distribution is a conjugate distribution to the multinomial.

Let $\vec{y} \sim \text{Multinomial}_K(n, \vec{\theta})$ and let $\vec{\theta} \sim \text{Dirichlet}_K(\vec{\alpha})$. Then we find the posterior distribution, $P(\vec{\theta}|\vec{y})$ is given by

$$P(\vec{\theta}|\vec{y}) = P(\vec{y}|\vec{\theta}) \cdot P(\vec{\theta}|\vec{\alpha}) \quad (3)$$

$$= \frac{n!}{\prod_{k=1}^K y_k!} \prod_{k=1}^K \theta_k^{y_k} \cdot \mathcal{B}(\vec{\alpha})^{-1} \prod_{k=1}^K \theta_k^{\alpha_k-1} \quad (4)$$

$$= \frac{n!}{\prod_{k=1}^K y_k!} \left(\prod_{k=1}^K \theta_k^{y_k} \right) \cdot \mathcal{B}(\vec{\alpha}) \left(\prod_{k=1}^K \theta_k^{\alpha_k-1} \right) \quad (5)$$

$$\propto \prod_{k=1}^K \theta_k^{y_k + \alpha_k - 1} \quad (6)$$

Thus with a Dirichlet prior on $\vec{\theta}$, the posterior, $\vec{\theta}|\vec{y}$ is distributed $\text{Dirichlet}_K(\vec{y} + \vec{\alpha})$.

Some other useful properties

One useful property is $\int_{\Omega} \prod_{k=1}^K \theta_k^{y_k + \alpha_k - 1} d\theta = \mathcal{B}(\vec{\alpha} + \vec{y})$

proof:

Since the posterior distribution of $\vec{\theta}$ is a probability distribution, we know it integrates to 1. Then,

$$\int_{\Omega} \mathcal{B}(\vec{\alpha} + \vec{y})^{-1} \prod_{k=1}^K \theta_k^{y_k + \alpha_k - 1} d\theta = 1 \quad (7)$$

$$\mathcal{B}(\vec{\alpha} + \vec{y})^{-1} \int_{\Omega} \prod_{k=1}^K \theta_k^{y_k + \alpha_k - 1} d\theta = 1 \quad (8)$$

$$\int_{\Omega} \prod_{k=1}^K \theta_k^{y_k + \alpha_k - 1} d\theta = \mathcal{B}(\vec{\alpha} + \vec{y}) \quad (9)$$

Another useful property is that the Dirichlet distribution is a member of the *exponential family* of distributions. Exponential family distributions can be written using the common form:

$$P(\vec{\theta}|\vec{\eta}) = h(\vec{\theta}) \exp \left\{ \vec{\eta}^T t(\vec{\theta}) - a(\vec{\eta}) \right\} \quad (10)$$

where

- η is called the *natural parameter*,
- $t(\vec{\theta})$ is a *sufficient statistic* for $\vec{\theta}$,
- $h(\vec{\theta})$ is the *underlying measure*,
- and $a(\eta)$ is the *log normalizer*.

The log normalizer is denoted as $a(\eta) = \log \left(\int_{\Omega} h(\vec{\theta}) \exp \left\{ \vec{\eta}^T t(\vec{\theta}) \right\} d\vec{\theta} \right)$.

It can be shown that the Dirichlet distribution is an exponential family distribution by exponentiating the log of the Dirichlet's PDF.

$$P(\vec{\theta}|\vec{\alpha}) = \mathcal{B}(\vec{\alpha})^{-1} \prod_{k=1}^K \theta_k^{\alpha_k - 1} \quad (11)$$

$$\log(P(\vec{\theta}|\vec{\alpha})) = \left(\sum_{k=1}^K (\alpha_k - 1) \log(\theta_k) \right) - \log(\mathcal{B}(\vec{\alpha})) \quad (12)$$

$$P(\vec{\theta}|\vec{\alpha}) = \exp \left\{ \left(\sum_{k=1}^K (\alpha_k - 1) \log(\theta_k) \right) - \log(\mathcal{B}(\vec{\alpha})) \right\} \quad (13)$$

Note that:

- The natural parameter is $\eta_k = \alpha_k - 1$.
- The sufficient statistic is $t(\theta_k) = \log(\theta_k)$
- The log normalizer is $a(\vec{\eta}) = \log(\mathcal{B}(\vec{\alpha}))$

And it turns out that the derivatives of the log normalizer are the moments of the sufficient statistic:

$$E_p[t(\vec{\theta})] = \frac{\partial a}{\partial \vec{\eta}^T} \quad (14)$$

Priors and Conjugacy

Latent Dirichlet Allocation

Blei originally gave two formulations for LDA and he did not state it in a pure Bayesian manner. The description below differs on both counts. We only consider the *smoothed* formulation of LDA and state it as a pure Bayesian model.

First, for clarity, let's define our indices:

- K topics, indexed by $k \in \{1, 2, \dots, K\}$
- D documents, indexed by $d \in \{1, 2, \dots, D\}$
- V unique words, indexed by $v \in \{1, 2, \dots, V\}$
- N_d words in the d -th document, indexed by $n \in \{1, 2, \dots, N_d\}$




Figure 1: LDA Plate Diagram

The model is stated as:

$$\vec{\phi}_k \sim \text{Dirichlet}_V(\vec{\beta}) \quad (15)$$

$$\vec{\theta}_d \sim \text{Dirichlet}_K(\vec{\alpha}) \quad (16)$$

$$\vec{z}_{d,n} \sim \text{Multinomial}_K(1, \vec{\theta}_d) \quad (17)$$

$$\vec{w}_{z_{d,n}} \sim \text{Multinomial}_V(1, \vec{\phi}_{z_{d,n}}) \quad (18)$$

Each topic vector, $\vec{\phi}_k$, is a row in a $K \times V$ matrix, Φ . Each document vector, $\vec{\theta}_d$, is a row in a $D \times K$ matrix, Θ . The random variable $\vec{z}_{d,n}$ represents the topic of the n -th word sampled for the d -th document. The random variable $\vec{w}_{z_{d,n}}$ represents the n -th word sampled for the d -th document, coming from the k -th topic as selected by $\vec{z}_{d,n}$. And these variables can be collectively represented as rows of matrices \mathbf{Z} and \mathbf{W} , respectively.

In the topic modeling literature, this model is often represented as a *plate diagram*. This type of diagram represents conditional dependencies with a directed graph. An example is in Figure 1.

The plate diagram does help make the posterior more obvious:

$$P(\mathbf{W}, \mathbf{Z}, \Theta, \Phi | \vec{\alpha}, \vec{\beta}) = \prod_{n=1}^{N_d} P(z_{d,n} | \theta_d) P(w_{d,n} | \phi_{z_{d,n}}) \cdot \prod_{d=1}^D P(\theta_d | \alpha) \cdot \prod_{k=1}^K P(\phi_k | \beta) \quad (19)$$

Collapsed Gibbs Sampling for LDA